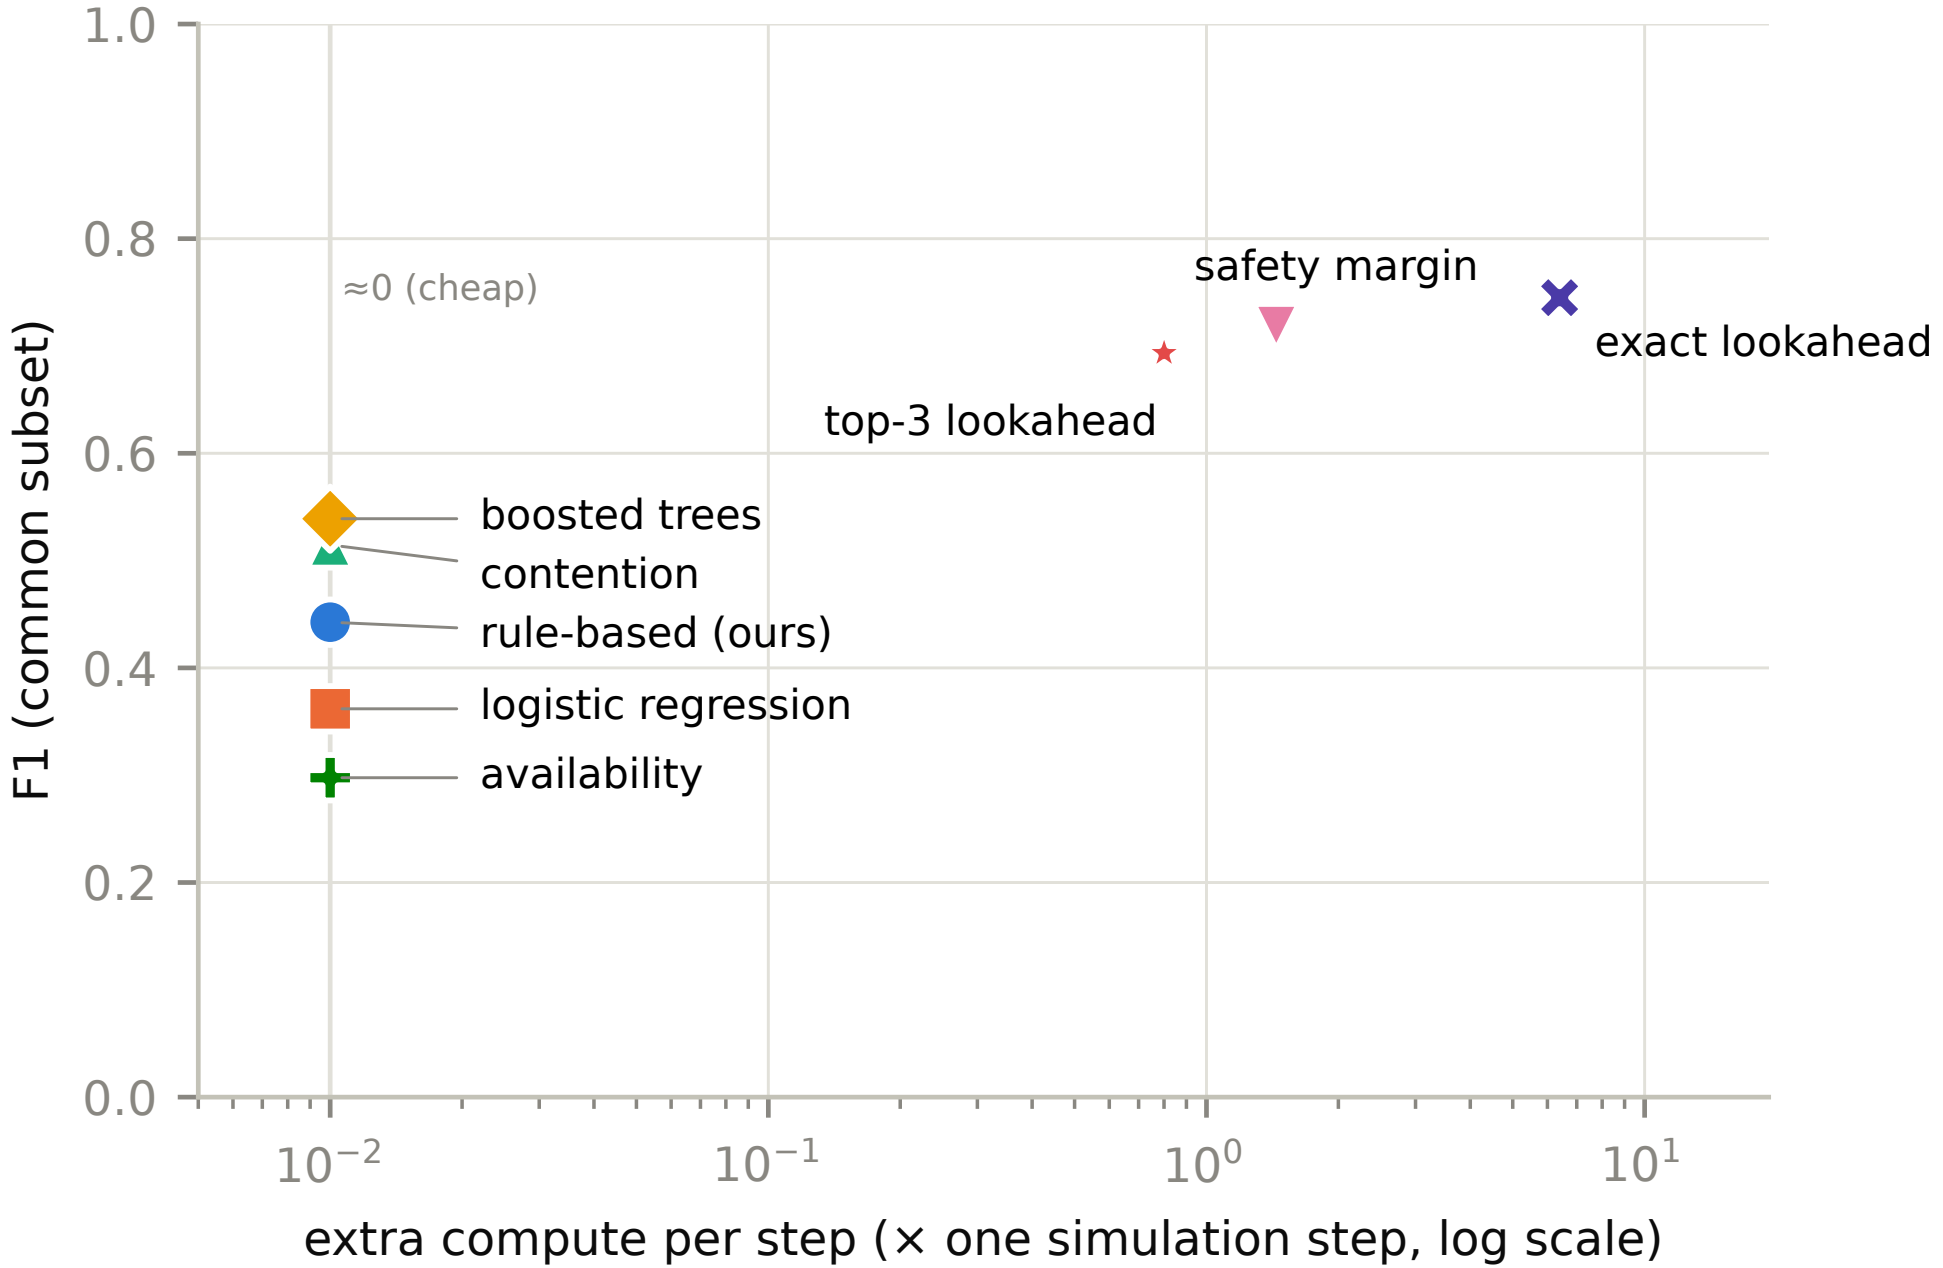
1.0
0.8
0.6
0.4
0.2
0.0
F1 (common subset)
≈0 (cheap)
safety margin
exact lookahead
top-3 lookahead
boosted trees
contention
rule-based (ours)
logistic regression
availability
10^{-2}
10^{-1}
10^{0}
10^{1}
extra compute per step (× one simulation step, log scale)